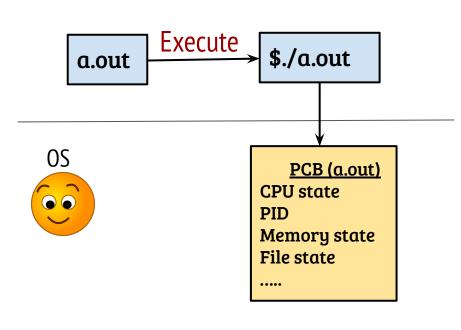
# CS330: Operating Systems

Process API: System calls

#### Recap: The process abstraction

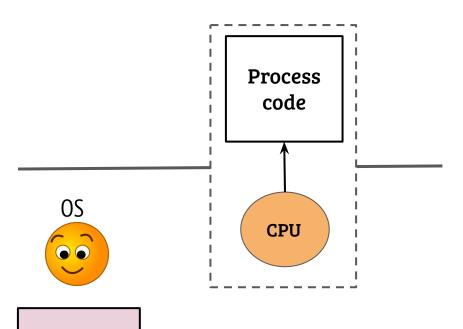
- The OS creates a process when we run an executable



- When we execute "a.out" on a shell a process control block (PCB) is created
- Does it raise some questions related to the exact working?

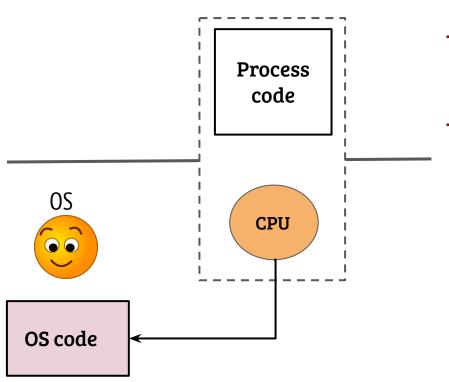
#### Process creation: What and How?

- How does OS come into action after typing "./a.out" in a shell?

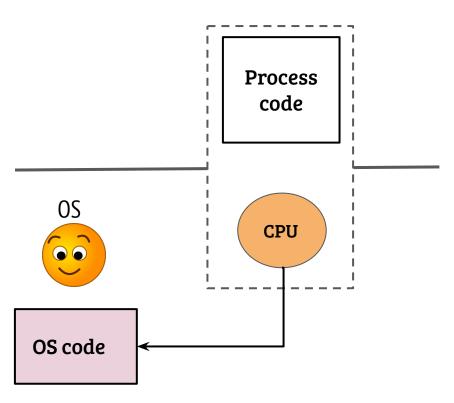


- CPU executing *user code* can invoke the *OS functions* using system calls

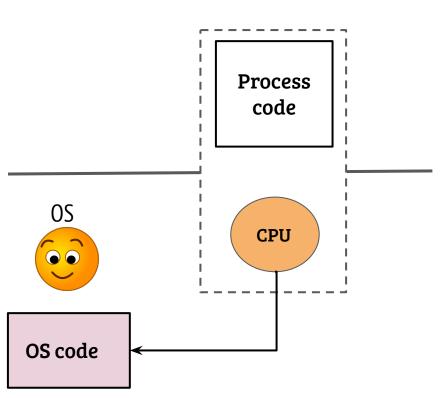
OS code



- CPU executing *user code* can invoke the *OS functions* using system calls
- The CPU executes the OS handler for the system call

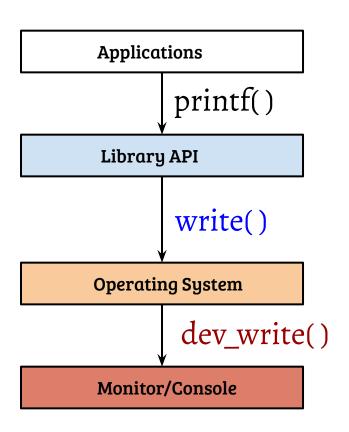


- CPU executing *user code* can invoke the *OS functions* using system calls
- The CPU executes the OS handler for the system call
- How system call is different from a function call?



- CPU executing *user code* can invoke the *OS functions* using system calls
- The CPU executes the OS handler for the system call
- How is system call different from a function call?
- Can be thought as an invocation of privileged functions (will revisit)

#### System calls and user libraries



- Most system calls are invoked through wrapper library functions
- However, all system calls can be invoked directly
  - For example, in Linux systems,
     syscall() wrapper can be used
     (Refer: man syscall)

# A simple system call: getpid()

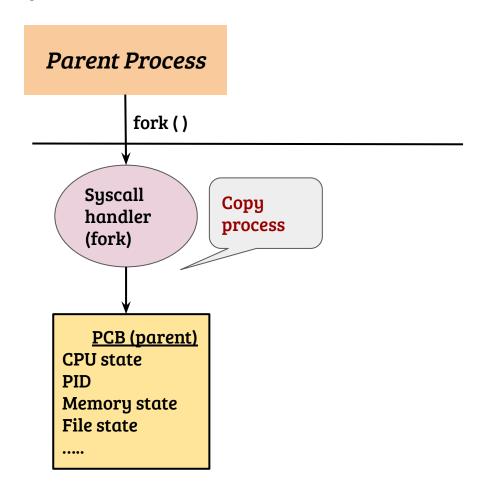
```
USER
                             pid_t getpid()
main()
                                PCB *current = get_current_process();
 printf("%d\n", getpid());
                                return (current → pid);
```

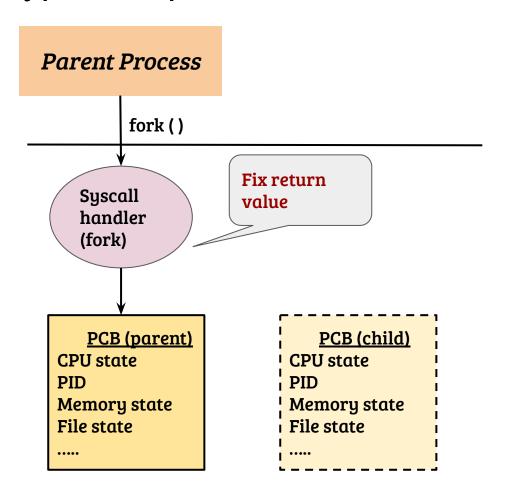
#### Process creation: What and How?

- How does OS come into action after typing "./a.out" in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?

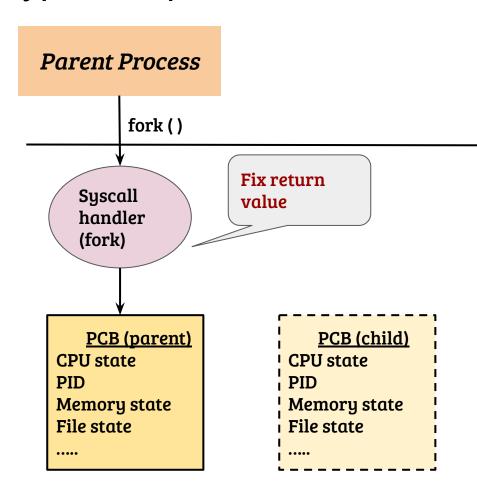
# Parent Process Fork() Child Process

- fork() system call is weird; not a typical "privileged" function call
- fork() creates a new process; a *duplicate* of calling process
- On success, fork
  - Returns PID of child process to the caller (parent)
  - Returns 0 to the child

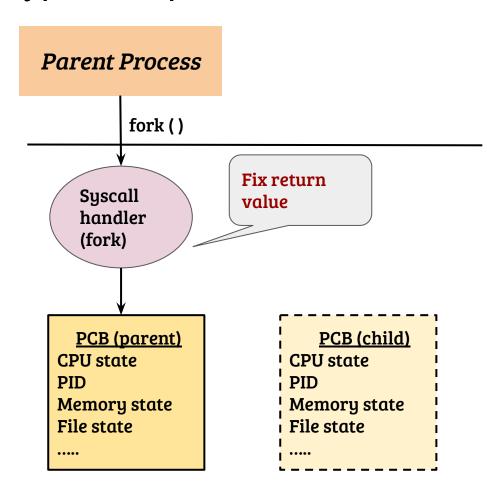




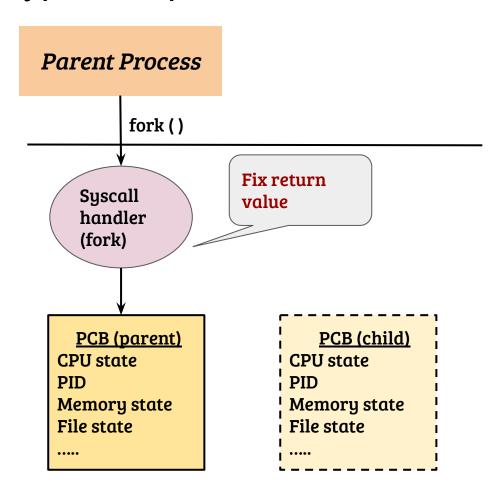
- Child should get '0' and parent gets PID of child as return value. How?



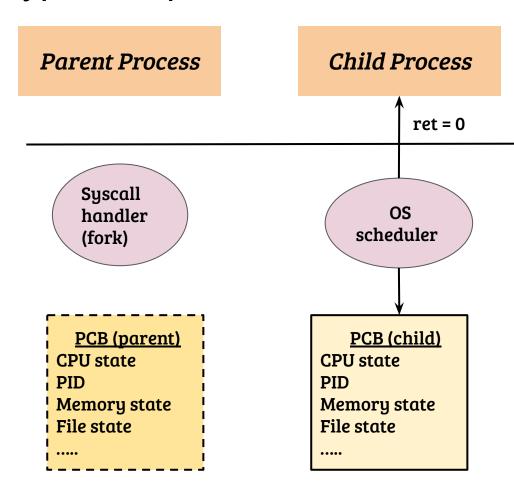
- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child



- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child
- When does child execute?



- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child
- When does child execute?
- When OS schedules the child process



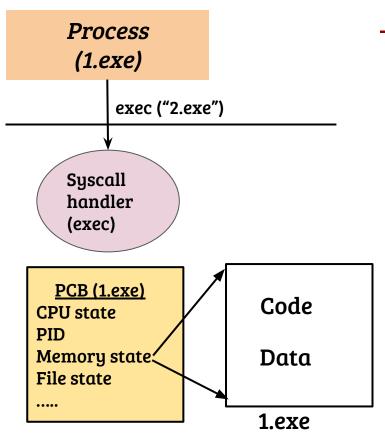
- PC is next instruction after fork() syscall, for both parent and child
- Child memory is an exact copy of parent
- Parent and child diverge from this point

#### Load a new binary - exec()



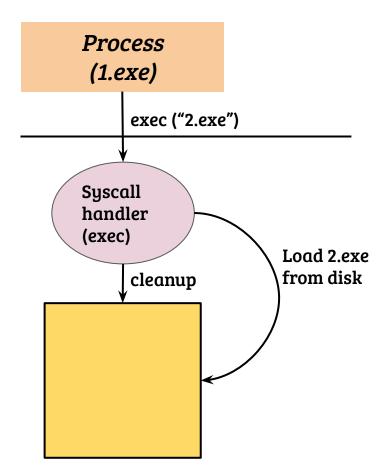
- Replace the calling process by a new executable
  - Code, data etc. are replaced by the new process
  - Usually, open files remain open

#### Typical implementation of exec



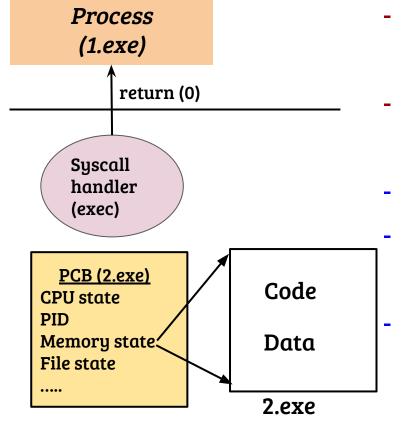
- The calling process commits self destruction! (almost)

#### Typical implementation of exec



- The calling process commits self destruction! (almost)
- The calling process is cleaned up and replaced by the new executable
- PID remains the same

#### Typical implementation of exec

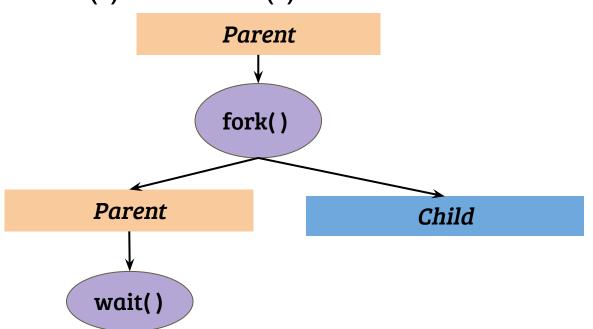


- The calling process commits self destruction! (almost)
- The calling process is cleaned up and replaced by the new executable
- PID remains the same
- On return, new executable starts execution
  - PC is loaded with the starting address of the newly loaded binary

#### Process creation: What and How?

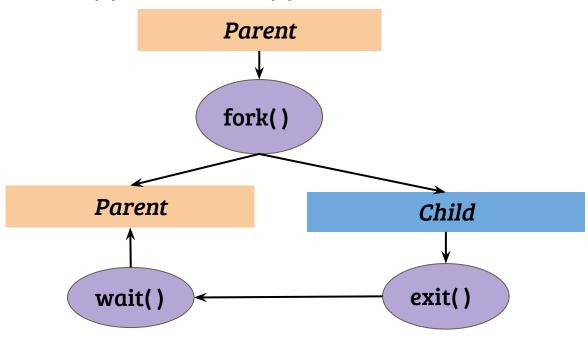
- How does OS come into action after typing "./a.out" in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?
- fork( ), exec ( ), wait( ) and exit( )
- Who invokes the system calls? In what order?

# wait() and exit()



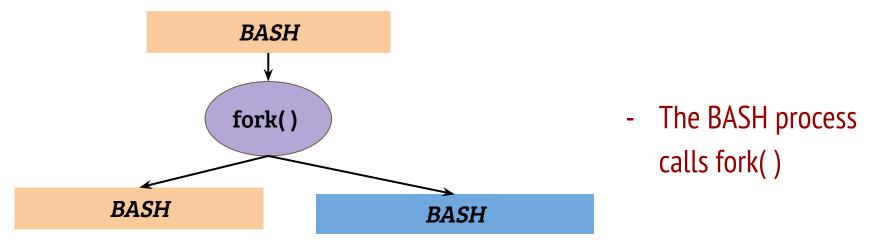
 The wait system call makes the parent wait for child process to exit

# wait( ) and exit( )

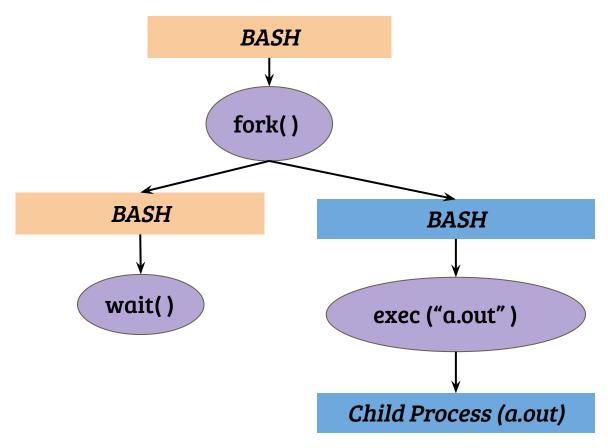


- The wait system call makes the parent wait for child process to exit
- On child exit(), the wait() system call returns in parent

#### Shell command line: fork + exec + wait

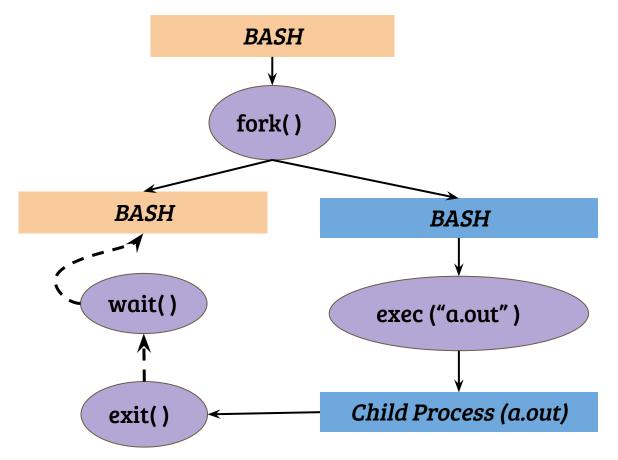


#### Shell command line: fork + exec + wait



- Parent process callswait() to wait for childto finish
- Child process invokes exec( )

#### Shell command line: fork + exec + wait

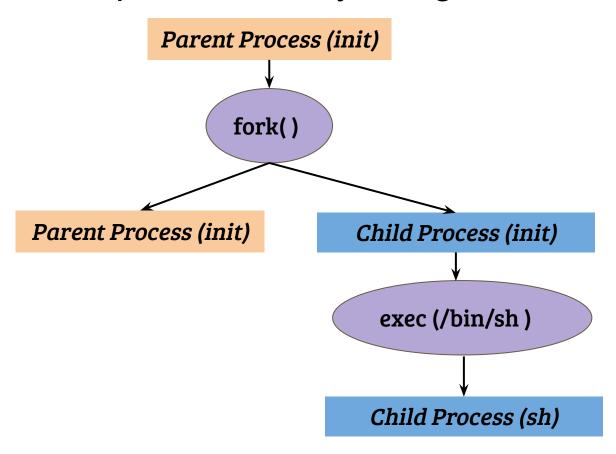


- When child exits, parent gets notified
- The BASH shell is ready for the next command at this point of time

#### Process creation: What and How?

- How does OS come into action after typing "./a.out" in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?
- fork( ), exec ( ), wait( ) and exit( )
- Who invokes the system calls? In what order?
- The shell process (bash process)
- What is the first user process?

#### Unix process family using fork + exec



- Fork and exec are used to create the process tree
- Commands: ps, pstree
- See the /proc directory in linux systems

#### Process creation: What and How?

- How does OS come into action after typing "./a.out" in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?
- fork(), exec(), wait() and exit()
- Who invokes the system calls?
- The shell process (bash process)
- What is the first user process?
- In Unix systems, it is called the *init* process
- Who creates and schedules the init process?