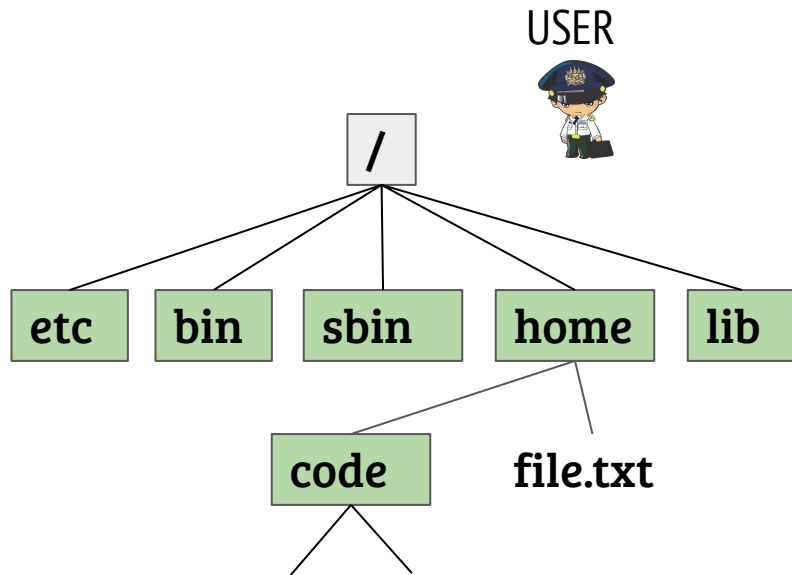


# CS330: Operating Systems

Files

# The file system



End-user wants see a nice tree view. Let me enable it through a simple system call APIs.



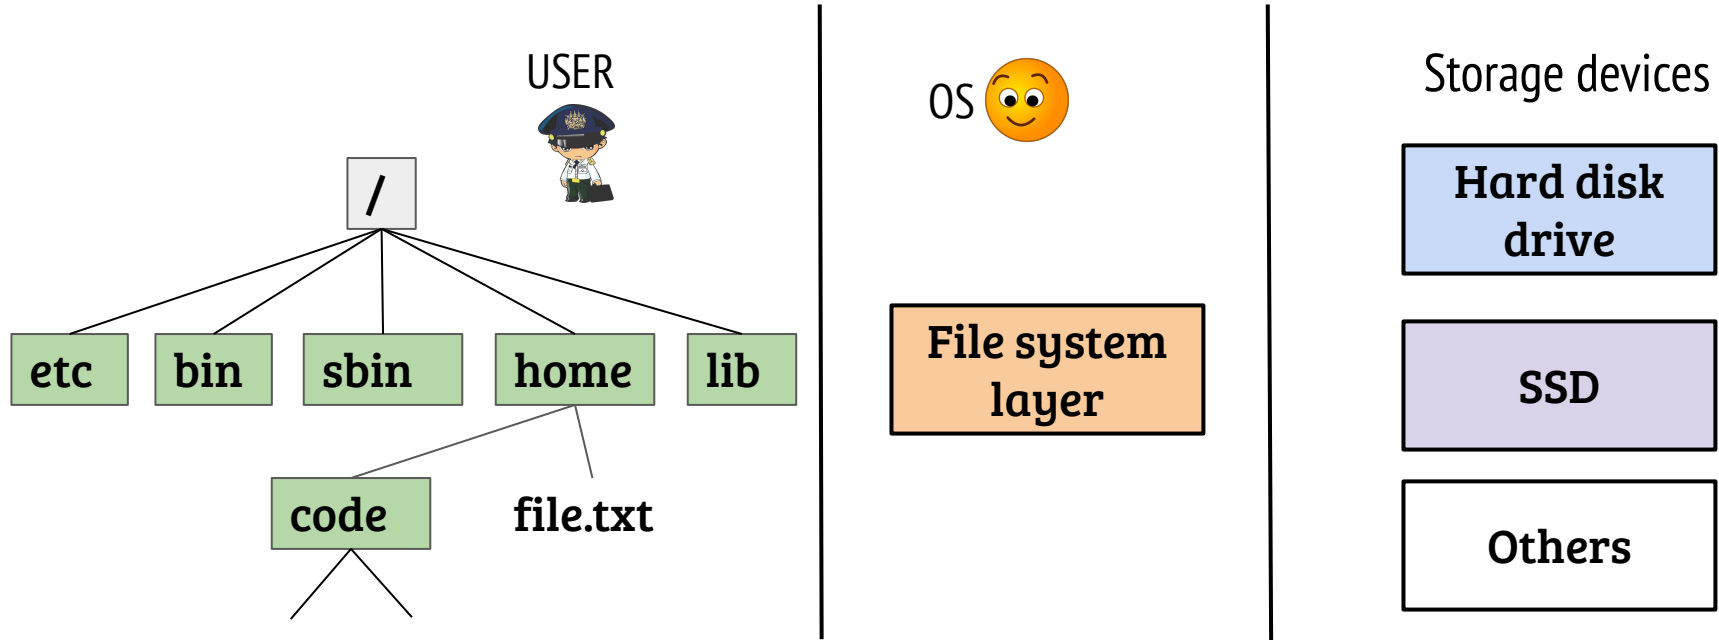
Storage devices

**Hard disk drive**

**SSD**

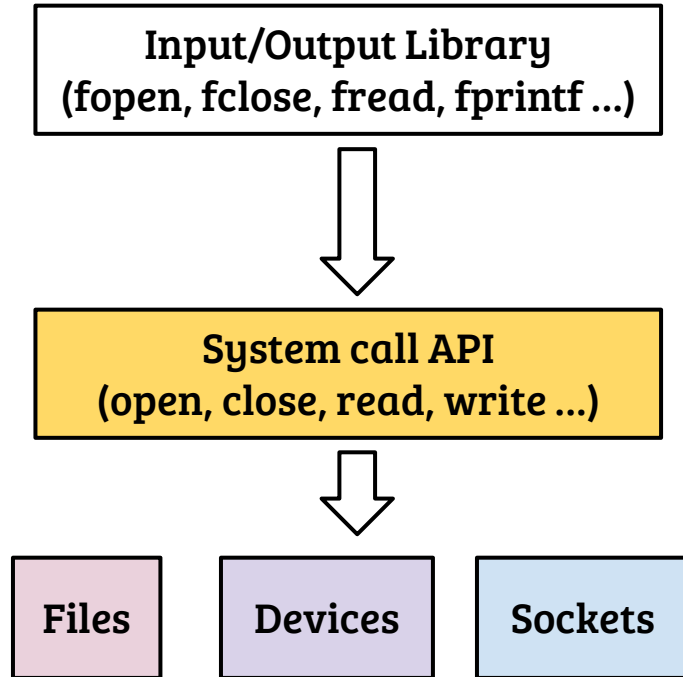
**Others**

# The file system



- File system is an important OS subsystem
  - Provides abstractions like files and directories
  - Hides the complexity of underlying storage devices

# File system interfacing



- Processes identify files through a file handle a.k.a. file descriptors
- In UNIX, the POSIX file API is used to access files, devices, sockets etc.
- What is the mapping between library functions and system calls?

open: getting a handle

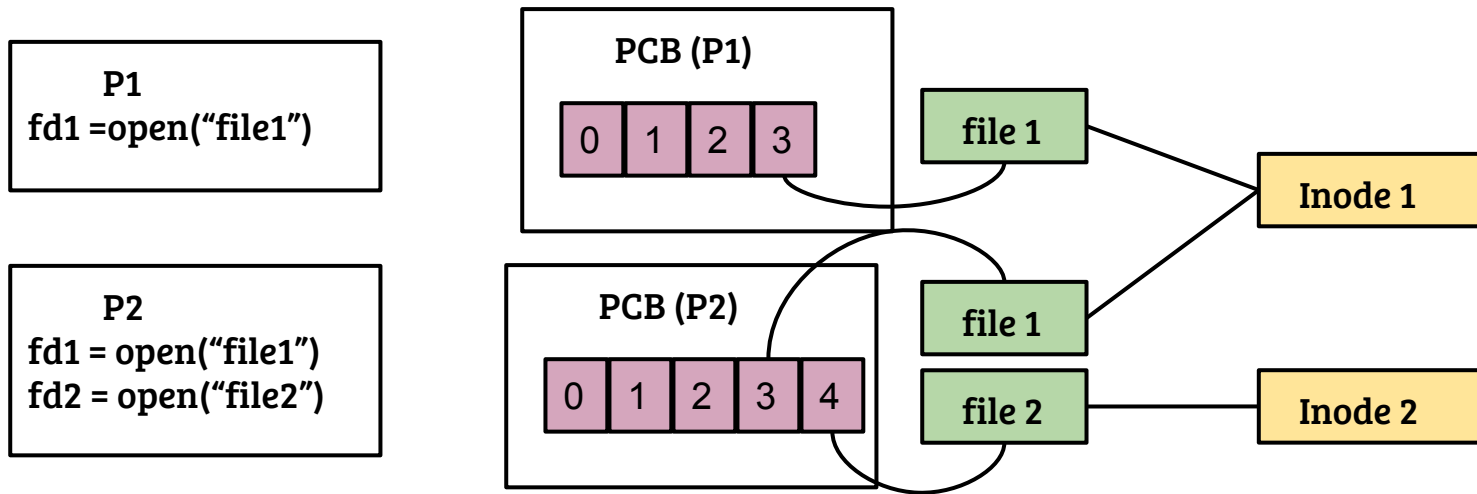
```
int open (char *path, int flags, mode_t mode)
```

# open: getting a handle

```
int open (char *path, int flags, mode_t mode)
```

- Access mode specified in flags : O\_RDONLY, O\_RDWR, O\_WRONLY
- Access permissions check performed by the OS
- On success, a file descriptor (integer) is returned
- If flags contain O\_CREAT, mode specifies the file creation mode
- Refer man page ("man 2 open")

# Process view of file



- Per-process file descriptor table with pointer to a “file” object
- file object → inode is many-to-one

# Process view of file



- What do file descriptors 0, 1 and 2 represent?
- What happens to the FD table and the file objects across `fork( )`?
  - What happens in `exec( )`?
- Can multiple FDs point to the same file object?



# Read and Write

```
ssize_t read (int fd, void *buf, size_t count);
```

- fd → file handle
- buf → user buffer as read destination
- count → #of bytes to read
- read ( ) returns #of bytes actually read, can be smaller than count

# Read and Write

```
ssize_t read (int fd, void *buf, size_t count);
```

- fd → file handle
- buf → user buffer as read destination
- count → #of bytes to read
- read ( ) returns #of bytes actually read, can be smaller than count

```
ssize_t write (int fd, void *buf, size_t count);
```

- Similar to read

# Process view of file



- What do file descriptors 0, 1 and 2 represent?
- 0  $\rightarrow$  STDIN, 1  $\rightarrow$  STDOUT and 2  $\rightarrow$  STDERR
- What happens to the FD table and the file objects across `fork( )`?
  - What happens in `exec( )`?
- Can multiple FDs point to the same file object?

# lseek

```
off_t lseek(int fd, off_t offset, int whence);
```

- fd → file handle
- offset → target offset
- whence → SEEK\_SET, SEEK\_CUR, SEEK\_END
- On success, returns offset from *the starting of the file*

# lseek

`off_t lseek(int fd, off_t offset, int whence);`

- `fd` → file handle
- `offset` → target offset
- `whence` → `SEEK_SET`, `SEEK_CUR`, `SEEK_END`
- On success, returns offset from *the starting of the file*
- Examples
  - `lseek(fd, SEEK_CUR, 100)` → forwards the file position by 100 bytes
  - `lseek(fd, SEEK_END, 0)` → file pos at EOF, returns the file size
  - `lseek(fd, SEEK_SET, 0)` → file pos at beginning of file

# File information (stat, fstat)

```
int stat(const char *path, struct stat *sbuf);
```

- Returns the information about file/dir in the argument `path`
- The information is filled up in structure called `stat`

# File information (stat, fstat)

```
int stat(const char *path, struct stat *sbuf);
```

- Returns the information about file/dir in the argument `path`
- The information is filled up in structure called `stat`

```
struct stat sbuf;
```

```
stat("/home/user/tmp.txt", &sbuf);
```

```
printf("inode = %d size = %ld\n", sbuf.st_ino, sbuf.st_size);
```

- Other useful fields in *struct stat* : `st_uid`, `st_mode` (Refer stat man page)

# Process view of file

PCB (P1)

- What do file descriptors 0, 1 and 2 represent?
- $0 \rightarrow \text{STDIN}$ ,  $1 \rightarrow \text{STDOUT}$  and  $2 \rightarrow \text{STDERR}$
- What happens to the FD table and the file objects across `fork( )`?
  - What happens in `exec( )`?
- The FD table is copied across `fork( )`  $\Rightarrow$  File objects are shared
- On `exec`, open files remain shared by default
- Can multiple FDs point to the same file object?



# Duplicate file handles (dup and dup2)

```
int dup(int oldfd);
```

- The dup() system call creates a “copy” of the file descriptor `oldfd`
- Returns the lowest-numbered unused descriptor as the new descriptor
- The old and new file descriptors represent the same file

# Duplicate file handles (dup and dup2)

```
int fd, dupfd;
```

```
fd = open("tmp.txt");
```

```
close(1);
```

```
dupfd = dup(fd);    //What will be the value of dupfd?
```

```
printf("Hello world\n"); // Where will be the output?
```

# Duplicate file handles (dup and dup2)

```
int fd, dupfd;
```

```
fd = open("tmp.txt");
```

```
close(1);
```

```
dupfd = dup(fd);    //What will be the value of dupfd?
```

```
printf("Hello world\n"); // Where will be the output?
```

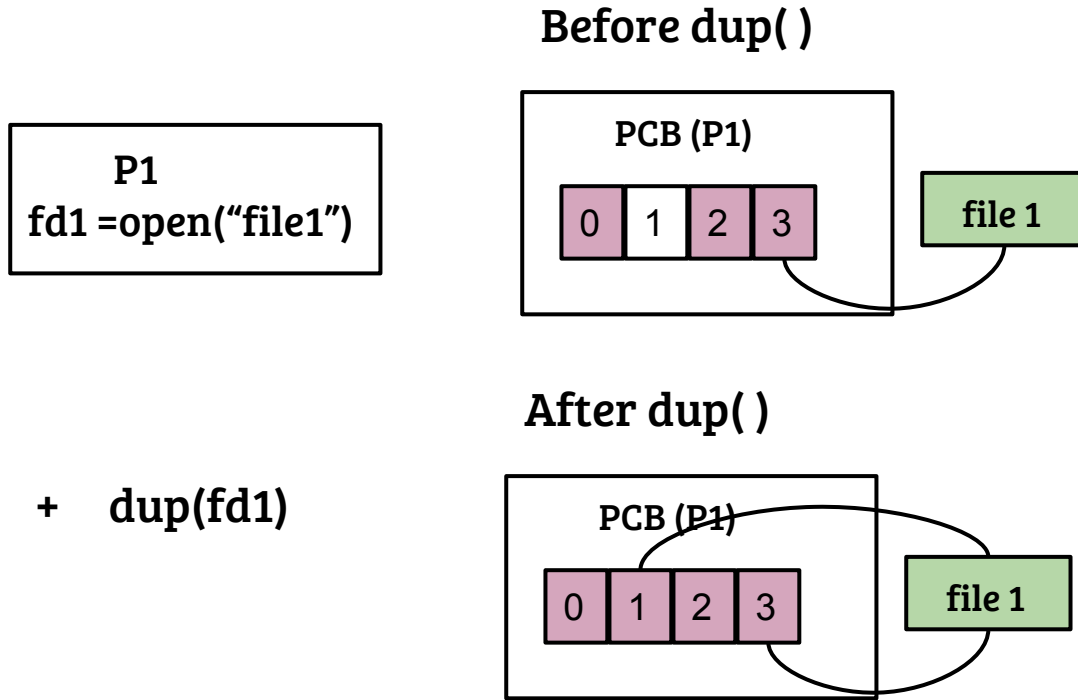
- Value of dupfd = 1 (assuming STDIN is open)
- "Hello world" will be written to tmp.txt file

# Duplicate file handles (dup and dup2)

```
int dup2(int oldfd, int newfd);
```

- Close `newfd` before duping the file descriptor `oldfd`
- `dup2 (fd, 1)` equivalent to
  - `close(1);`
  - `dup(fd);`

# Duplicate file handles (dup and dup2)



- Lowest numbered unused fd (i.e., 1) is used (Assume STDOUT is closed before)
- Duplicate descriptors share the same file state
- Closing one file descriptor *does not* close the file

# Use of dup: shell redirection

- Example: `ls > tmp.txt`
- How implemented?

# Use of dup: shell redirection

- Example: `ls > tmp.txt`
- How implemented?

```
fd = open ("tmp.txt")
```

```
close( 1); close(2);    // close STDOUT and STDERR
```

```
dup(fd); dup(fd)        // 1→ fd, 2 → fd
```

```
exec(ls )
```

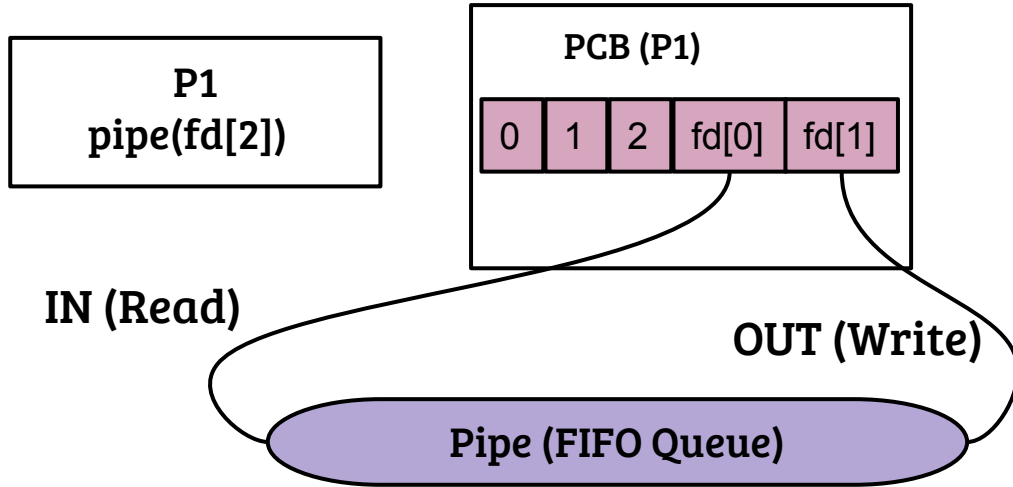
# Process view of file

PCB (P1)

- What do file descriptors 0, 1 and 2 represent?
- 0  $\rightarrow$  STDIN, 1  $\rightarrow$  STDOUT and 2  $\rightarrow$  STDERR
- What happens to the FD table and the file objects across `fork( )`?
  - What happens in `exec( )`?
- The FD table is copied across `fork( )`  $\Rightarrow$  File objects are shared
- On `exec`, open files remain shared by default
- Can multiple FDs point to the same file object?
- Yes, `duped` FDs share the same file object (within a process)

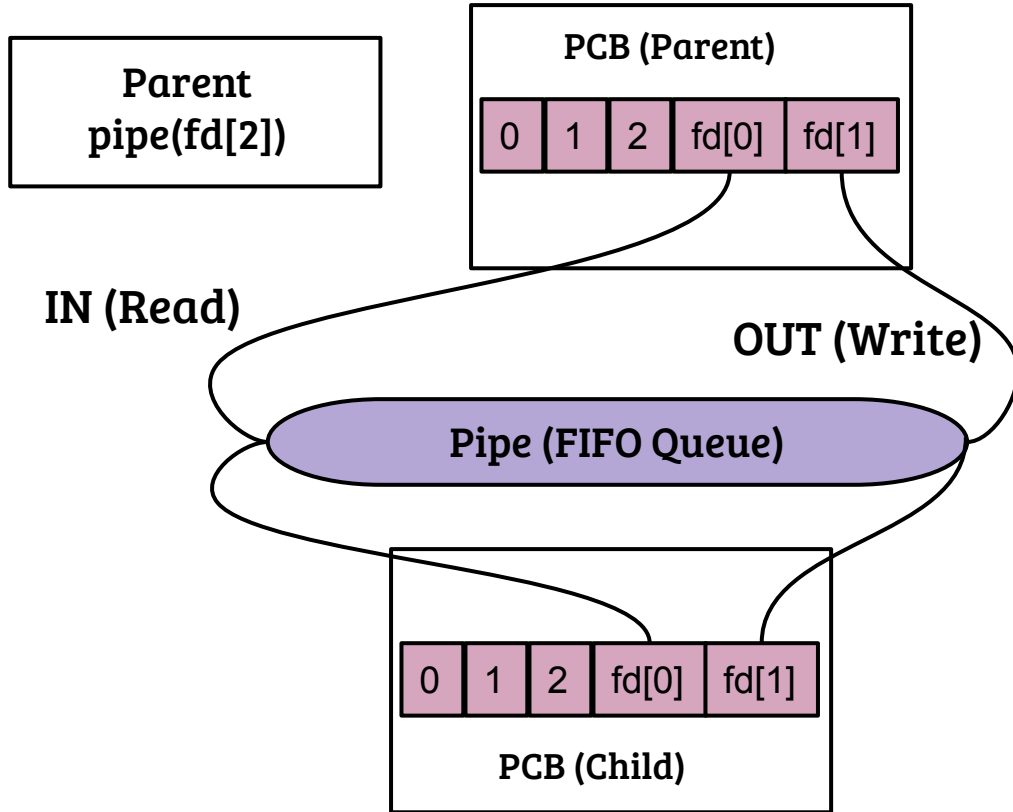


# UNIX pipe( ) system call



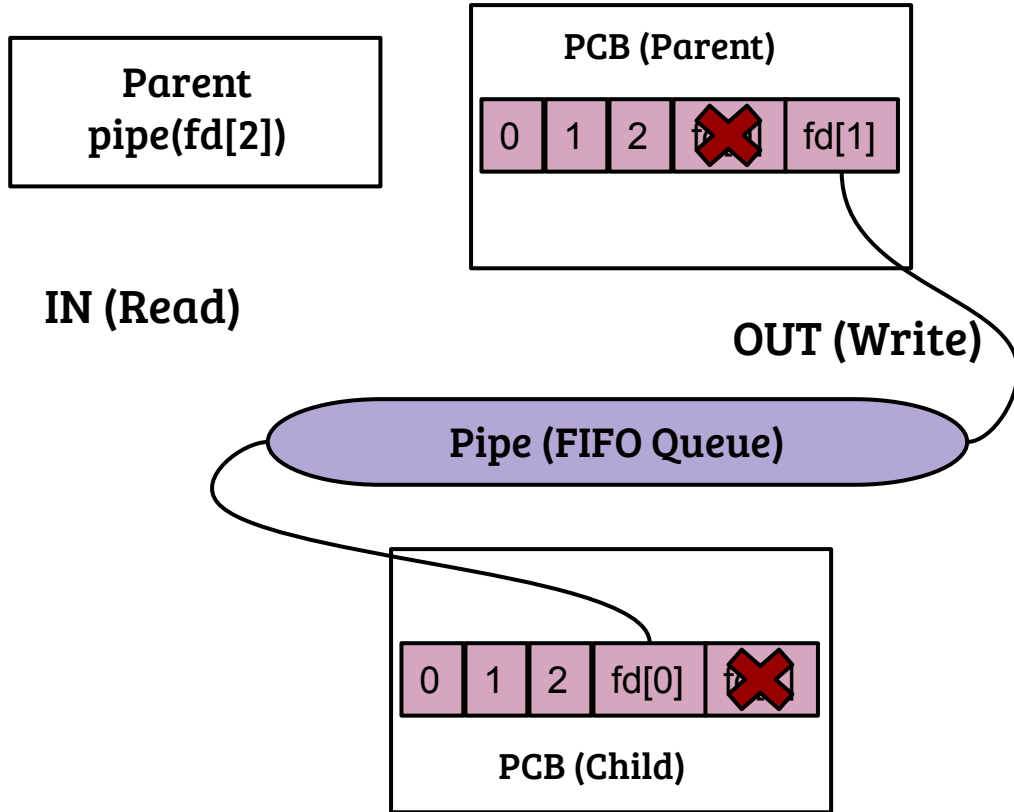
- `pipe( )` takes array of two FDs as input
- `fd[0]` is the read end of the pipe
- `fd[1]` is the write end of the pipe
- Implemented as a FIFO queue in OS

# UNIX pipe( ) with fork( )



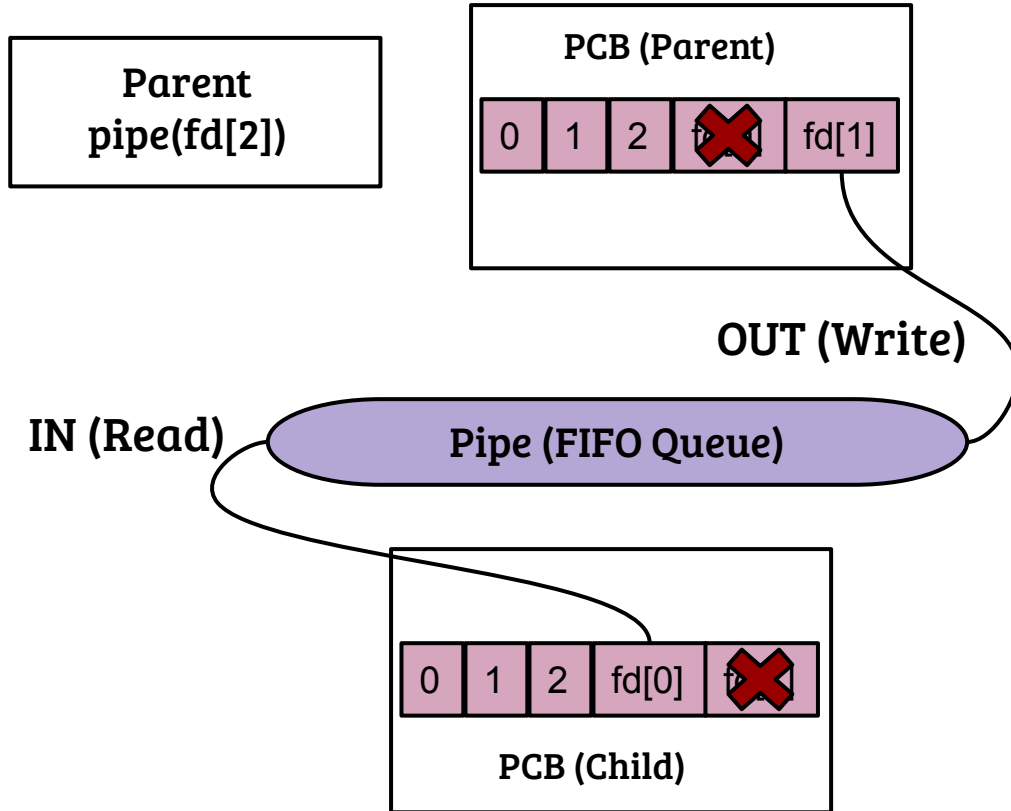
- `fork( )` duplicates the file descriptors
- At this point, both the parent and the child processes can read/write to the pipe

# UNIX pipe( ) with fork( )



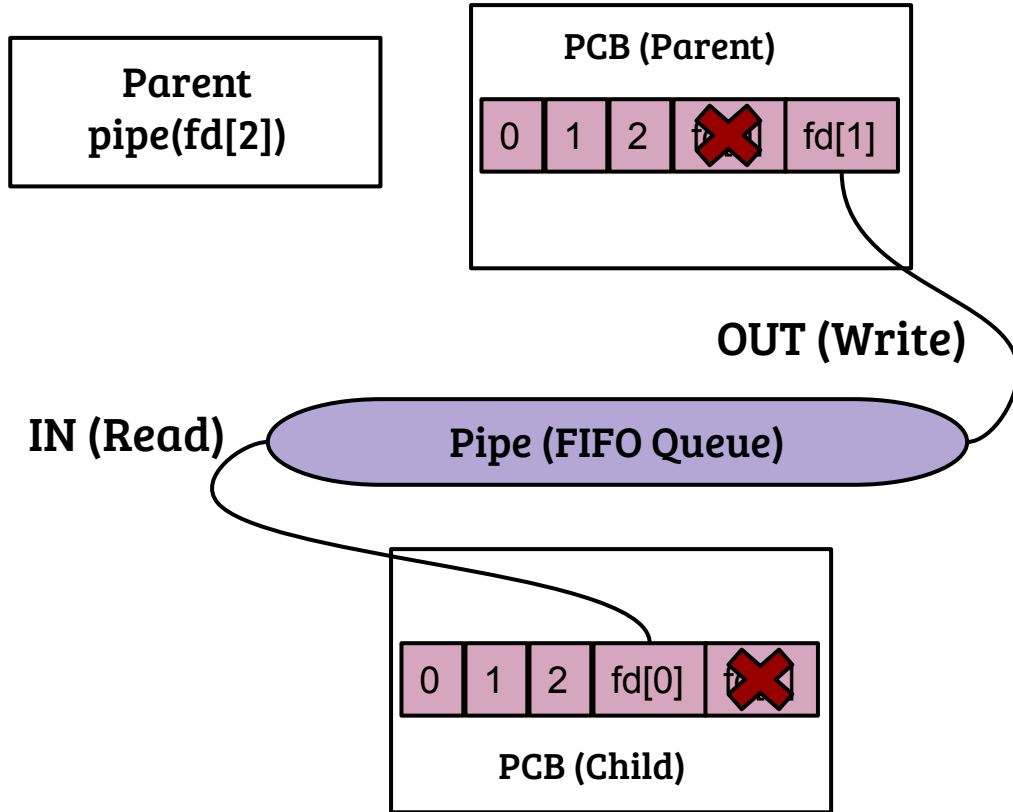
- fork( ) duplicates the file descriptors
- close( ) one end of the pipe, both in child and parent
- Result
  - A queue between parent and child

# Shell piping : ls | wc -l



- `pipe( )` followed by `fork( )`
- Parent: `exec( "ls" )` after making `STDOUT`  $\rightarrow$  out fd of the pipe (using `dup`)

# Shell piping : ls | wc -l



- `pipe( )` followed by `fork( )`
- Parent: `exec( "ls" )` after making `STDOUT`  $\rightarrow$  out fd of the pipe (using `dup`)
- Child: `exec( "wc" )` after closing `STDIN` and duping in fd of pipe
- Result: input of "wc" is connected to output of "ls"