

CS330: Operating Systems

Locking

Recap: Synchronization and locking

- Locking is necessary when multiple contexts access shared resources
- Example: Multiple threads, multiple OS execution contexts
- Efficiency of lock and unlock operations
- Hardware-assisted lock implementations are used for efficiency
- Lock acquisition delay vs. wasted CPU cycles
- Use waiting locks and spinlocks depending on the requirement
- Fairness of the locking scheme
- Contending threads should not starve for the lock (infinitely)

Recap: Synchronization and locking

- Locking is necessary when multiple contexts access shared resources
- Example: Multiple threads, multiple OS execution contexts
- Efficiency of lock and unlock operations
- Hardware-assisted lock implementations are used for efficiency
- Lock acquisition delay vs. wasted CPU cycles
- Use waiting locks and spinlocks depending on the requirement
- Fairness of the locking scheme
- Contending threads should not starve for the lock

Agenda: Spinlocks, Semaphore and mutex (waiting locks)

Spinlock: Buggy attempt

1. `lock_t *L; // Initial value = 0` - Does this implementation work?
2. `lock(L)`
3. `{`
4. `while(*L);`
5. `*L = 1;`
6. `}`
7. `unlock(L)`
8. `{`
9. `*L = 0;`
10. `}`

Spinlock: Buggy attempt

1. `lock_t *L; // Initial value = 0` - Does this implementation work?
2. `lock(L)` - No, it does not ensure *mutual exclusion*
3. `{` - Why?
4. `while(*L);`
5. `*L = 1;`
6. `}`
7. `unlock(L)`
8. `{`
9. `*L = 0;`
10. `}`

Spinlock: Buggy attempt

```
1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.     while(*L);
5.     *L = 1;
6. }
7. unlock(L)
8. {
9.     *L = 0;
10. }
```

- Does this implementation work?
- No, it does not ensure *mutual exclusion*
- Why?
 - Single core: Context switch between line #4 and line #5
 - Multicore: Two cores exiting the while loop by reading lock = 0

Spinlock: Buggy attempt

1. `lock_t *L; // Initial value = 0`
 2. `lock(L)`
 3. `{`
 4. `while(*L);`
 5. `*L = 1;`
 6. `}`
 7. `unlock(L)`
 8. `{`
 9. `*L = 0;`
 10. `}`
- Does this implementation work?
 - No, it does not ensure *mutual exclusion*
 - Why?
 - Single core: Context switch between line #4 and line #5
 - Multicore: Two cores exiting the while loop by reading `lock = 0`
 - Core issue: Compare and swap has to happen atomically!

Spinlock using atomic exchange

```
1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.     while(atomic_xchg(*L, 1));
5. }
6. unlock(L)
7. {
8.     *lock = 0;
9. }
```

- Atomic exchange: exchange the value of memory and register atomically
- `atomic_xchg (int *PTR, int val)` returns the value at PTR before exchange
- Ensures mutual exclusion if “val” is stored on a register
- No fairness guarantees

Spinlock using XCHG on X86

```
lock(lock_t *L)
{
    asm volatile(
        "mov $1, %%rax;"
        "loop: xchg %%rax, (%%rdi);"
        "cmp $0, %%rax;"
        "jne loop;"
        ::: "memory" );
}

unlock(int *L) { *L = 0;}
```

- $XCHG\ R, M \Rightarrow$ Exchange value of register R and value at memory address M
- RDI register contains the lock argument
- Exercise: Visualize a context switch between any two instructions and analyse the correctness

Spinlock using compare and swap

```
1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.     while( CAS(*L, 0, 1) );
5. }
6. unlock(L)
7. {
8.     *lock = 0;
9. }
```

- Atomic compare and swap: perform the condition check and swap atomically
- CAS (int **PTR*, int *cmpval*, int *newval*) sets the value of *PTR* to *newval* if *cmpval* is equal to value at *PTR*. Returns 0 on successful exchange
- No fairness guarantees!

CAS on X86: cmpxchg

cmpxchg source[Reg] destination [Mem/Reg]

Implicit registers : rax and flags

1. if rax == [destination]
2. then
3. flags[ZF] = 1
4. [destination] = source
5. else
6. flags[ZF] = 0
7. rax = [destination]

- “cmpxchg” is not atomic in X86, should be used with a “lock” prefix

Spinlock using CMPXCHG on X86

```
lock(lock_t *L)
{
asm volatile(
    "mov $1, %%rcx;"
    "loop: xor %%rax, %%rax;"
    "lock cmpxchg %%rcx, (%%rdi);"
    "jnz loop;"
    ::: "rcx", "rax", "memory");
}

unlock(lock_t *L) { *L = 0;}
```

- Value of RAX (=0) is compared against value at address in register RDI and exchanged with RCX (=1), if they are equal
- Exercise: Visualize a context switch between any two instructions and analyse the correctness

Load Linked (LL) and Store conditional (SC)

- LoadLinked (R, M)
 - Like a normal load, it loads R with value of M
 - Additionally, the hardware keeps track of future stores to M
- StoreConditional (R, M)
 - Stores the value of R to M if no stores happened to M after the execution of LL instruction (after execution, R = 1)
 - Otherwise, store is not performed (after execution R=0)
- Supported in RISC architectures like mips, risc-v etc.

Spinlock using LL and LC

```
lock_t *L; //initial value = 0    lock:  LL R1, (R2); //R2 = lock address
lock(lock_t *L)                  BNEQZ R1, lock;
{                                ADDUI R1, R0, #1; //R1 = 1
    while(LoadLinked(L) ||      SC R1, (R2)
        !StoreConditional(L, 1));
    BEQZ R1, lock
}
unlock(lock_t *L) { *L = 0;}
```

- Efficient as the hardware avoids memory traffic for unsuccessful lock acquire attempts
- Context switch between LL and SC results in SC to fail

Spinlocks: reducing wasted cycles

- Spinning for locks can introduce significant CPU overheads and increase energy consumption
- How to reduce spinning in spinlocks?

Spinlocks: reducing wasted cycles

- Spinning for locks can introduce significant CPU overheads and increase energy consumption
- How to reduce spinning in spinlocks?
- Strategy: Back-off after every failure, exponential back-off used mostly

```
lock( lock_t *L) {  
    u64 backoff = 0;  
    while(LoadLinked(L) || !StoreConditional(L, 1)){  
        if(backoff < 63) ++backoff;  
        pause(1 << backoff); // Hint to processor  
    }  
}
```


Fairness in spinlocks

- Spinlock implementations discussed so far are not fair,
 - no bounded waiting
- To ensure fairness, some notion of ordering is required
- What if the threads are granted the lock in the order of their arrival to the lock contention loop?
 - A single lock variable may not be sufficient
 - Example solution: Ticket spinlocks

Atomic fetch and add (xadd on X86)

xadd R, M

TmpReg T = R + [M]

R = [M]

[M] = T

- Example: M = 100; RAX = 200
- After executing “lock xadd %RAX, M”, value of RAX = 100, M = 300
- Require lock prefix to be atomic

Ticket spinlocks (OSTEP Fig. 28.7)

```
struct lock_t{
    long ticket;
    long turn;
};


void init_lock (struct lock_t *L){
    L → ticket = 0; L → turn = 0;
}

void unlock(struct lock_t *L){
    L → turn++;
}
```

```
void lock(struct lock_t *L){
    long myturn = xadd(&L → ticket, 1);
    while(myturn != L → turn)
        pause(myturn - L → turn);
}
```


- Example: Order of arrival: T1 T2 T3
- T1 (in CS) : myturn = 0, L = {1, 0}
- T2: myturn = 1, L = {2, 0}
- T3: myturn = 2, L = {3, 0}
- T1 unlocks, L = {3, 1}. T2 enters CS

Ticket spinlock




- Local variable “myturn” is equivalent to the order of arrival
- If a thread is in CS \Rightarrow Local Turn must be same as “Turn”
- Threads waiting = Ticket - Turn - 1

Ticket spinlock



- Value of turn incremented on lock release
- Thread which arrived just after the current thread enters the CS
- When a new thread arrives, it gets the lock after the other threads ahead of the new thread acquire and release the lock

Ticket spinlock



- Ticket spinlock guarantees bounded waiting
- If N threads are contending for the lock and execution of the CS consumes T cycles, then $\text{bound} = N * T$ (assuming negligible context switch overhead)

Ticket spinlock (with yield)

```
void lock(struct lock_t *L){  
    long myturn = xadd(&L → ticket, 1);  
    while(myturn != L → turn)  
        sched_yield( );  
}
```

- Why spin if the thread's turn is yet to come?
- Yield the CPU and allow the thread with ticket (or other non contending threads)
- Further optimization
 - Allow the thread with “myturn” value one more than “L → turn” to continue spinning