

# CS330: Operating Systems

Read-write locks, classical problems

# Recap: Spinlocks and semaphores

- Lock implementation using atomic exchange, compare-and-swap, LL-SC and atomic add
- Ticket spinlocks provide fairness in locking, example implementation with atomic-add (xadd)
- Software-only lock implementation (Peterson's solution)
- Semaphore: counting semaphore, binary semaphore (mutex)

Agenda: Read-write locks, producer consumer problem, concurrency bugs

# Reader-writer locks

- Allows *multiple readers* or *a single writer* to enter the CS
- Example: Insert, delete and lookup operations on a search tree

# Reader-writer locks

- Allows *multiple readers* or *a single writer* to enter the CS
- Example: Insert, delete and lookup operations on a search tree

```
struct BST{  
    struct node *root;  
    rwlock_t *lock;  
};
```

```
struct node{  
    item_t item;  
    struct node *left;  
    struct node *right;  
};
```

```
void insert(BST *t, item_t item);  
void lookup(BST *t, item_t item);
```

# Reader-writer locks

- Allows *multiple readers* or *a single writer* to enter the CS
- Example: Insert, delete and lookup operations on a search tree

```
struct BST{  
    struct node *root;  
    rwlock_t *lock;  
};
```

```
struct node{  
    item_t item;  
    struct node *left;  
    struct node *right;  
};
```

```
void insert(BST *t, item_t item);  
void lookup(BST *t, item_t item);
```

- If multiple threads call lookup( ), they may traverse the tree in parallel

# Implementation of read-write locks

```
struct rwlock_t{  
    Lock read_lock;  
    Lock write_lock;  
    int num_readers;  
}
```

```
init_lock(rwlock_t *rL)  
{  
    init_lock(&rL → read_lock);  
    init_lock(&rL → write_lock);  
    rL → num_readers = 0;  
}
```

# Implementation of read-write locks (writers)

```
struct rwlock_t{  
    Lock read_lock;  
    Lock write_lock;  
    int num_readers;  
}
```

```
void write_lock(rwlock_t *rL)  
{  
    lock(&rL → write_lock);  
}
```

```
init_lock(rwlock_t *rL)  
{  
    init_lock(&rL → read_lock);  
    init_lock(&rL → write_lock);  
    rL → num_readers = 0;  
}
```

```
void write_unlock(rwlock_t *rL)  
{  
    unlock(&rL → write_lock);  
}
```

- Write lock behavior is same as the typical lock, only one thread allowed to acquire the lock

# Implementation of read-write locks (readers)

```
struct rwlock_t{  
    Lock read_lock;  
    Lock write_lock;  
    int num_readers;  
}
```

```
void read_lock(rwlock_t *rL)  
{  
    lock(&rL → read_lock);  
    rL → num_readers++;  
    if(rL → num_readers == 1)  
        lock(&rL → write_lock);  
    unlock(&rL → read_lock);  
}
```

```
void read_unlock(rwlock_t *rL)  
{  
    lock(&rL → read_lock);  
    rL → num_readers--;  
    if(rL → num_readers == 0)  
        unlock(&rL → write_lock);  
    unlock(&rL → read_lock);  
}
```



# Implementation of read-write locks (readers)

```
struct rwlock_t{  
    Lock read_lock;  
    Lock write_lock;  
    int num_readers;  
}
```

```
void read_lock(rwlock_t *rL)  
{  
    lock(&rL → read_lock);  
    rL → num_readers++;  
    if(rL → num_readers == 1)  
        lock(&rL → write_lock);  
    unlock(&rL → read_lock);  
}
```

- The first reader acquires the write lock prevents writers to acquire lock
- The last reader releases the write lock to allow writers

```
void read_unlock(rwlock_t *rL)  
{  
    lock(&rL → read_lock);  
    rL → num_readers--;  
    if(rL → num_readers == 0)  
        unlock(&rL → write_lock);  
    unlock(&rL → read_lock);  
}
```

# Producer-consumer problem

```
DoProducerWork(){
```


```
while(1){
```

```
    item_t item = prod_p();
```

```
    produce(item);
```

```
}
```

```
}
```



```
DoConsumerWork(){
```

```
while(1){
```

```
    item_t item = consume();
```

```
    cons_p(item);
```

```
}
```

```
}
```

- A buffer of size N, one or more producers and consumers
- Producer produces an element into the buffer (after processing)
- Consumer extracts an element from the buffer and processes it
- Example: A multithreaded web server, network protocol layers etc.
- How to solve this problem using semaphores?

# Buggy #1

```
item_t A[n], pctr=0, cctr = 0;  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used);  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    sem_post(&empty);  
    return item;  
}
```

- This solution does not work. What is the issue?

# Buggy #1

```
item_t A[n], pctr=0, cctr = 0;  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used);  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    sem_post(&empty);  
    return item;  
}
```

- This solution does not work. What is the issue?
- The counters (pctr and cctr) are not protected, can cause race conditions

# Buggy #2

```
item_t A[n], pctr=0, cctr = 0; lock_t *L = init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    lock(L); sem_wait(&empty);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    sem_post(&used); unlock(L);  
}
```

```
item_t consume() {  
    lock(L); sem_wait(&used);  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    sem_post(&empty); unlock(L);  
    return item;  
}
```

- What is the problem?

# Buggy #2

```
item_t A[n], pctr=0, cctr = 0; lock_t *L = init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    lock(L); sem_wait(&empty);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    sem_post(&used); unlock(L);  
}
```

```
item_t consume( ) {  
    lock(L); sem_wait(&used);  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    sem_post(&empty); unlock(L);  
    return item;  
}
```

- What is the problem?
- Consider empty = 0 and producer has taken lock before the consumer. This results in a deadlock, consumer waits for L and producer for empty

# A working solution

```
item_t A[n], pctr=0, cctr = 0; lock_t *L = init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty); lock(L);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    unlock(L); sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used); lock(L)  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    unlock(L); sem_post(&empty);  
    return item;  
}
```

- The solution is deadlock free and ensures correct synchronization, but very much serialized (inside produce and consume)
- What if we use separate locks for producer and consumer?

# Solution with separate mutexes

```
item_t A[n], pctr=0, cctr = 0; lock_t *P = init_lock(), *C=init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty); lock(P);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    unlock(P); sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used); lock(C)  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    unlock(C); sem_post(&empty);  
    return item;  
}
```

- Does this solution work?
- Homework: Assume that item is a large object and copy of item takes long time. How can we perform the copy operation without holding the lock?